

# Advanced Shenanigans

Yufeng Huang

Associate Professor of Marketing, Simon Business School

August 18, 2022

## Dictionary



# she·nan·i·gans

/SHĕˈnænəɡənz/

*noun*

INFORMAL

plural noun: **shenanigans**; noun: **shenanigan**

secret or dishonest activity or maneuvering.

"widespread financial shenanigans had ruined the fortunes of many"

- silly or high-spirited behavior; mischief.

### Origin

mid 19th century: of unknown origin.

Translate shenanigans to

Choose language



Use over time for: shenanigans



Show less

From Oxford

Feedback

# What I mean is: let's spend a bit of time on algorithms

- ▶ But these are advanced applications
  - ▶ optional for the course; but I think they complement what we know nicely
- ▶ Specifically, we'll start with manually solving for a nonlinear equation (using a loop)
  - ▶ variation of this example: optimization of a (differentiable) function
  - ▶ application: profit maximization
  - ▶ goal: a) loops/functions, b) understanding how the optimizer works, c) understanding that different algorithms will lead to drastic differences in efficiency
  - ▶ why optimization per se? You'll find it useful across many applications of statistical models or profit functions of a company

# Solving a nonlinear equation

# Solving a non-linear equation

- ▶ As the most basic problem, consider solving for a nonlinear equation of the form

$$f(x) = 0$$

where  $f$  is a differentiable function (i.e. its derivative exists)

- ▶ For this class, I will introduce the standard Newton-Raphson algorithm, i.e. the following recurve formula will arrive at the solution when  $t \rightarrow \infty$ :

$$x_{t+1} = x_t - \frac{f(x_t)}{f'(x_t)}$$

where  $f'$  is the derivative of the function  $f$

## But how to find the derivative of $f$ ?

- ▶ A simple way to take numerical derivatives is to find the slope of  $f(x)$  at a very small change of  $x$ , i.e. to calculate

$$f'(x) = \frac{f(x + \Delta x) - f(x)}{\Delta x}$$

- ▶ For example, let's take the derivative for

$$f(x) = x^2 - 4$$

which some of us know (from some math classes) is  
 $f'(x) = 2x$

```
# let the function f be
f <- function(x) x^2 - 4

# then the derivative at x = 1 is
(f(1 + 0.001) - f(1)) / 0.001

## [1] 2.001

# which we can confirm if we know how to do the math
```

# Example 1

- Now, let's find **one** solution for

$$x^2 - 4 = 0$$

```
# let the function f be
f <- function(x) x^2 - 4

# then the solution should be to continue the iteration
x_new <- 0 # starting point
for (i in 1:100) {      # run the iteration 100 times
  # replace previous x_old
  x_old <- x_new

  # take derivative (gradient) of f
  fx <- f(x_old)
  dfx <- (fx - f(x_old - 0.001)) / 0.001

  # evaluate the next x
  x_new <- x_old - fx / dfx
}

# print results
x_new

## [1] -2
```

# Comments

- ▶ Good starting point, but:
  - ▶ the 100 iterations seem unnecessary or not enough, depending on how difficult it is to find the solution
  - ▶ better is to decide whether to continue the loop using conditions about whether we “*converged*” at a solution
- ▶ Natural choice is a while loop:
  - ▶ with an explicit condition to continue/stop the iteration
  - ▶ a typical stopping criteria is to stop the iteration if changes in  $x$  is smaller than some tolerance
  - ▶ more sophisticated is to require both convergence in  $x$  and convergence in  $f(x)$ , but we'll keep it simple here



# Example 1 modified

```
# a better choice is a while loop
x_new <- 0 # starting point
tol_X <- 1E-6 # stopping criteria
x_old <- 1 # placeholder, you'll see why we need this

while (abs(x_new - x_old) > tol_X) { # continue if diff in x is large
  # replace previous x_old
  x_old <- x_new

  # take derivative (gradient) of f
  fx <- f(x_old)
  dfx <- (fx - f(x_old - 0.001)) / 0.001

  # evaluate the next x
  x_new <- x_old - fx / dfx
}

# print results
x_new

## [1] -2
```


## Example 2

- Now, let's find **one** solution for

$$5x^3 - 7x^2 - 40x + 100 = 0$$

but this time do it using a function

```
# let the function f be
f <- function(x) {
  5 * x^3 - 7 * x^2 - 40 * x + 100
}
# show the function in a figure
range_x <- seq(-4, 1, by = 0.05)
plot(range_x, f(range_x), type = 'l', xlab = 'x', ylab = 'f(x)')
```



## Example 2

```
# write the routine into a function:
solve_f <- function(par, fn, tolX = 1E-6, gr = NULL) {

  # initialize
  x_new <- par
  x_old <- par + 1

  # loop
  while (abs(x_new - x_old) > tolX) {

    # replace previous x_old
    x_old <- x_new
    fx <- fn(x_old)

    # take derivative (gradient) of f (address this later)
    if (is.null(gr)) {
      # if gradient does not exist
      dfx <- (fx - fn(x_old - 0.001)) / 0.001
    } else {
      # otherwise, gradient is the gradient function evaluated at x_old
      dfx <- gr(x_old)
    }

    # evaluate the next x
    x_new <- x_old - fx / dfx

  }

  return(list(par = x_new, f = fn(x_new)))
}
```

## Example 2

```
# call the solve_f function with standard arguments
solve_f(par = 0, fn = f)

## $par
## [1] -3.151719
##
## $f
## [1] -2.960164e-09

# however, now that we can specify the gradient argument gr, let's see what it does
# recall that we can numerically take the gradient as something like (fx - f(x_old - 0.001)) / 0.001
# however, we can explicitly take the derivative as use it as the gradient
# for function f here, the derivative is  $5 * 3 * x^2 - 7 * 2 * x - 40$ 
g_f <- function(x) {
  5 * 3 * x^2 - 7 * 2 * x - 40
}

# now, call solve_f with the gradient argument
solve_f(par = 0, fn = f, gr = g_f)

## $par
## [1] -3.151719
##
## $f
## [1] -4.263256e-14

# why am I covering this?
# optim and other solvers also use the gr argument, this is the way to provide it
```

# Optimization

## Example 3: optimization


- ▶ At this point one asks: why are we learning how to solve nonlinear equations?
- ▶ Answer: this is the fundamental to many applications
- ▶ Example of such an application: we will code up our own optimizer to replace optim
  - ▶ well, we'll stick to a one-dimensional problem for simplicity
- ▶ To be specific: we want to maximize the function

$$\text{objfun}(x) = -0.02x^4 + 0.08x^3 - 0.3x^2 + 0.8x + 5$$

how do we do this by hand?

## Example 3

```
# let function objfun be
objfun <- function(x) {
  -0.02 * x^4 + 0.08 * x^3 - 0.3 * x^2 + 0.8 * x + 5
}
# examine the function
range_x <- seq(-2, 4, by = 0.05)
plot(range_x, objfun(range_x), type = 'l', xlab = 'x', ylab = 'f(x)')
```



## Example 3

- ▶ How do we optimize this? Note that for differentiable functions, a necessary condition for the maximum (or minimum) is where the derivative is zero
- ▶ So we can use the previously-defined `solve_f()` function to find the point where derivative is zero

```
# take numerical derivative
numd_objfun <- function(x) {
  (objfun(x + 0.0001) - objfun(x)) / 0.0001
}

# find the point where the derivative is zero
x1 <- solve_f(par = 0, fn = numd_objfun)
x1

## $par
## [1] 1.859048
##
## $f
## [1] 8.881784e-12

# what is the function value?
objfun(x1$par)

## [1] 5.725532

# compare against standard package
optim(0, function(x) -objfun(x))$par

## [1] 1.859375
```



## Example 3

- ▶ What we've seen is essentially the `optim()` function in the simplest form (or the Excel solver, etc.)
  - ▶ but there are many ways to improve it
  - ▶ for example, using analytical derivatives instead of numerical ones

```
# take analytical derivative (i.e. ones we derive using math)
d_objfun <- function(x) {
  -0.02 * 4 * x^3 + 0.08 * 3 * x^2 - 0.3 * 2 * x + 0.8
}

# find the point where the analytical derivative is zero
x2 <- solve_f(par = 0, fn = d_objfun)
x2$par # slightly different

## [1] 1.859098

# can even supply the second derivative
# which will be useful in solving for df(x) = 0
d2_objfun <- function(x) {
  -0.02 * 4 * 3 * x^2 + 0.08 * 3 * 2 * x - 0.3 * 2
}
x3 <- solve_f(par = 0, fn = d_objfun, gr = d2_objfun)
x3$par

## [1] 1.859098
```

# Summary

- ▶ Two points to learn from this section
  - ▶ how do we implement an algorithm using a while/for loop?
  - ▶ how does an optimizer work (in a simple way)?
- ▶ Next: real application
  - ▶ how do we maximize profit by finding the right quantity?
- ▶ Other applications?
  - ▶ how do we estimate a linear regression (or any linear and nonlinear models)?

# **Profit maximization in a competitive market (aka Samsung vs HTC redux)**

## Example 4: profit maximization and competition

- ▶ Recall that in the Samsung vs HTC case, we understood how to solve for the equilibrium in the following loop:
  - ▶ for each month, observe the market condition, define the game
  - ▶ for each combination of quantities  $q_H$  and  $q_S$ , determine the market price and profits
    - ▶ holding Samsung's quantity, HTC chooses  $q_H$  to maximize its profit
    - ▶ holding HTC's quantity, Samsung chooses  $q_S$  to maximize its profit
    - ▶ and so on and so forth...
- ▶ Looks cumbersome, right?
- ▶ Let's try to crack this problem in a simpler way

## Example 4: the old method

```
# take a look at the old method
#   first part of market.simulator.R

# ===== Step 1: construct profit matrix =====
pos.qty <- 1:nq

# initialize profit
profit_S <- array(NA, c(nq, nq))
profit_H <- array(NA, c(nq, nq))

# double for-loop to get the profit matrix
#   note: don't have /100 any more
for (q_S in pos.qty) {
  for (q_H in pos.qty) {
    price <- b1*income + b2*log(mktsize) - b3*(q_S + q_H) # more people: price decrease not as fast w
    profit_S[q_S, q_H] <- q_S * (price - mc_S)
    profit_H[q_S, q_H] <- q_H * (price - mc_H)
  }
}
```

## Example 4: the old method

```
# second part of market.simulator.R

# ===== Step 2: while loop to find best response / equilibrium quantity =====
# NOTE: This is what your market research team tells you

# initialize
i <- 300
j <- 100
change <- 1000

# keep iterating until no one wants to choose a different quantity
while (change > 0) {

  # 1. store previous values
  i_old <- i
  j_old <- j

  # 2. Samsung's response
  i <- which.max(profit_S[, j])

  # 3. HTC's response
  j <- which.max(profit_H[i, ])

  # 4. calculate changes from prev value
  change <- abs(i - i_old) + abs(j - j_old)

}

# quantities
q_S <- pos.qty[i]
q_H <- pos.qty[j]

# price
price <- b1*income + b2*log(mktsize) - b3*(q_S + q_H)
```

## Example 4: application of the old method

```
# reads data
df <- read.table('mobile_phone.csv', header = T, stringsAsFactors = F)

# parameters (say you know these)
nq <- 500
b1 <- 250
b2 <- 150
b3 <- 0.3

# generate 3 columns as the predicted quantity/price
df$quantity_S_pred <- NA
df$quantity_H_pred <- NA
df$price_pred <- NA

# run through all 72 periods (past and future)
for (t in 1:72) {
  # assign some values
  income <- df$income[t]
  mktsize <- df$mktsize[t]
  mc_S <- df$mc_S[t]
  mc_H <- df$mc_H[t]

  # run the market simulator every time period
  source('market.simulator.R')


  # collect results
  df$quantity_S_pred[t] <- q_S # recall that reacted.quantity is HTC - Samsung
  df$quantity_H_pred[t] <- q_H
  df$price_pred[t] <- price
}
```

```

# can compare predictions
df$t <- (df$year - min(df$year))*12 + df$month
plot(df$t, df$quantity_S_pred,
     type = 'l', lty = 2, xlab = "months since 201601", ylab = "quantity (dash = predicted)",
     main = "predicted/actual quantity by Samsung")
points(df$t, df$quantity_S, type = 'l')

```

**predicted/actual quantity by Samsung**





# However, how much time do we waste?

```
# make the loop into a function so that we can evaluate this
test1 <- function() {
  for (t in 1:72) {
    # assign some values
    income <- df$income[t]
    mktsize <- df$mktsize[t]
    mc_S <- df$mc_S[t]
    mc_H <- df$mc_H[t]

    # run the market simulator every time period
    source('market.simulator.R')

    # collect results
    df$quantity_S_pred[t] <- q_S # recall that reacted.quantity is HTC - Samsung
    df$quantity_H_pred[t] <- q_H
    df$price_pred[t] <- price
  }
}

library(rbenchmark)
benchmark(test1(), replications = 1)

##           test replications elapsed relative user.self sys.self user.child sys.child
## 1 test1()           1      5.28           1      5.17      0.1           NA           NA
```

# Thoughts

- ▶ In the above algorithm
  - ▶ have to calculate profits in every combinations of  $q_H$  and  $q_S$
  - ▶ and loop over 72 months
  - ▶ what would happen if  $q_H$  and  $q_S$  do not only have 500 possible values, but 50,000?
- ▶ We observe that maybe we do not have to exhaust all possible quantity combinations
  - ▶ e.g. some quantities result in low profits and will never be chosen
  - ▶ which leads to the next algorithm (plus the code will be more elegant)

# Improved algorithm (1)

```
# re-define the price and profit functions
find_price <- function(q1, q2, income, mktsize) {
  price <- b1*income + b2*log(mktsize) - b3*(q1 + q2)    # q1 and q2 interchangeable
  return(price)
}
find_profit <- function(q, q_opp, mc, income, mktsize) {
  profit <- q * (find_price(q, q_opp, income, mktsize) - mc)
  return(profit)
}
```

# Improved algorithm (2)

```
# first, setup the variables

# focus on month 1
income <- df$income[1]
mktsize <- df$mktsize[1]
mc_S <- df$mc_S[1]
mc_H <- df$mc_H[1]

# setup the initial variables
q_S_old <- 0
q_H_old <- 0
q_S <- 300
q_H <- 100
```

## Improved algorithm (3)

```
# for a given time period, simply iteratively optimize profit
#   until neither firm wants to deviate from a given quantity

while (abs(q_S - q_S_old) > 1 | abs(q_H - q_H_old) > 1) {

  # update the _old variables
  q_S_old <- q_S
  q_H_old <- q_H

  # Samsung maximizes profit
  q_S <- optim(par = q_S_old, fn = function(q) {
    -find_profit(q, q_H, mc_S, income, mktsize)
  })$par

  # HTC maximizes profit
  q_H <- optim(par = q_H_old, fn = function(q) {
    -find_profit(q, q_S, mc_H, income, mktsize)
  })$par

}

# examine results
q_S
## [1] 404.5878

q_H
## [1] 304.3925
```

```
# can now write the above into a function (that can go across many time periods)
```

```
solve_equilibrium <- function(t) {  
  
  # initialize  
  income <- df$income[t]  
  mktsize <- df$mktsize[t]  
  mc_S <- df$mc_S[t]  
  mc_H <- df$mc_H[t]  
  q_S_old <- 0  
  q_H_old <- 0  
  q_S <- 300  
  q_H <- 100  
  
  # loop  
  while (abs(q_S - q_S_old) > 1 | abs(q_H - q_H_old) > 1) {  
    # update the _old variables  
    q_S_old <- q_S  
    q_H_old <- q_H  
    # Samsung maximizes profit  
    q_S <- optim(par = q_S_old, fn = function(q) {  
      -find_profit(q, q_H, mc_S, income, mktsize)  
    })$par  
    # HTC maximizes profit  
    q_H <- optim(par = q_H_old, fn = function(q) {  
      -find_profit(q, q_S, mc_H, income, mktsize)  
    })$par  
  }  
  
  # return  
  return(c(q_S = q_S, q_H = q_H))  
}
```

```
solve_equilibrium(1)
```

```
##      q_S      q_H  
## 404.5878 304.3925
```

```
test2 <- function() {
  for (t in 1:72) {

    # solve for equilibrium
    res <- solve_equilibrium(t)

    # collect results
    df$quantity_S_pred[t] <- res["q_S"]
    df$quantity_H_pred[t] <- res["q_H"]

  }
}

# benchmark
benchmark(test1(), test2(), replications = 1)

##           test replications elapsed relative user.self sys.self user.child sys.child
## 1 test1()           1      5.18   43.167         5.00    0.17         NA         NA
## 2 test2()           1      0.12    1.000         0.12    0.00         NA         NA
```

- ▶ HUGE speed gain, because we have taken away much of the wasteful calculations!
- ▶ But, can we do even better??

## Pushing it even further...

- ▶ But, can we do even better??
- ▶ Note that instead of finding the maximum profit for each firm, holding the other firm's strategy fixed...
  - ▶ ... we can find such point where the derivative of profit is **zero** for each firm, holding the other firm's strategy fixed
  - ▶ of course, relies on a little bit of math, but the derivative of profit is still not very difficult: profit is

$$\text{profit}(q) = q \cdot (p(q) - mc)$$

then the derivative of it is

$$\begin{aligned}\text{profit}'(q) &= (p(q) - mc) + q \cdot p'(q) \\ &= (p(q) - mc) - b_3 \cdot q\end{aligned}$$

where the second line follows the linear inverse demand function



```

# construct the derivative function
d_profit <- function(q, q_opp, mc, income, mktsize, b3) {
  (find_price(q, q_opp, income, mktsize) - mc) - b3 * q
}

# now, can solve for zero-derivatives
solve_equilibrium_2 <- function(t) {
  # initialize
  income <- df$income[t]
  mktsize <- df$mktsize[t]
  mc_S <- df$mc_S[t]
  mc_H <- df$mc_H[t]
  q_S_old <- 0; q_H_old <- 0; q_S <- 300; q_H <- 100

  # loop
  while (abs(q_S - q_S_old) > 1 | abs(q_H - q_H_old) > 1) {
    # update the _old variables
    q_S_old <- q_S
    q_H_old <- q_H
    # Samsung finds q_S where derivative of profit is zero
    q_S <- solve_f(par = q_S_old, fn = function(q) {
      d_profit(q, q_H, mc_S, income, mktsize, b3)
    })$par
    # HTC finds q_H where derivative of profit is zero
    q_H <- solve_f(par = q_H_old, fn = function(q) {
      d_profit(q, q_S, mc_H, income, mktsize, b3)
    })$par
  }

  # return
  return(c(q_S = q_S, q_H = q_H))
}

solve_equilibrium_2(1)

##      q_S      q_H
## 404.5658 304.4161

```

## Putting it all together

- ▶ Now, you might ask: what's the point of defining a new `d_profit` function just to use our `solve_f()`?
  - ▶ well, if you think that's the only modification we're going to do, you're under-estimating the potential for the derivative function!
- ▶ The true purpose: we observe that the equilibrium is reached where for both firms, the derivative of profit is zero, i.e.

$$\text{profit}'_S(q_S) | q_H = 0$$

for Samsung and for HTC,

$$\text{profit}'_H(q_H) | q_S = 0$$

- ▶ This gives us a system of equations that will completely get around the while loop, and instead replace it with a system of two equations (which we'll solve using a package)

```

# MODIFY the derivative function such that
d_profit_vec <- function(q_S, q_H, mc_S, mc_H, income, mktsize, b3) {
  p <- find_price(q_S, q_H, income, mktsize)
  c((p - mc_S) - b3 * q_S, (p - mc_H) - b3 * q_H)
}

# load a nonlinear equation solver
library(nleqslv)

# define this procedure on one period
solve_equilibrium_3 <- function(t) {
  # initialize
  income <- df$income[t]
  mktsize <- df$mktsize[t]
  mc_S <- df$mc_S[t]
  mc_H <- df$mc_H[t]

  # solve the system using fsolve()
  res <- nleqslv(f = function(q) d_profit_vec(q[1], q[2], mc_S, mc_H, income, mktsize, b3),
    x = c(300, 100))
  # return
  return(c(q_S = res$x[1], q_H = res$x[2]))
}

# try it on period 1!
res <- solve_equilibrium_3(1)
res[1:2]

##      q_S      q_H
## 404.466 304.466

```

```

# loop over periods 1-72
test3 <- function() {
  for (t in 1:72) {
    res <- solve_equilibrium_3(t)
  }
  return(res)
}

# benchmark
benchmark(test1(), test2(), test3(), replications = 10) ###

##      test replications elapsed relative user.self sys.self user.child sys.child
## 1 test1()          10    52.50   1050.0     51.25     1.2         NA         NA
## 2 test2()          10     1.17    23.4      1.17     0.0         NA         NA
## 3 test3()          10     0.05     1.0      0.05     0.0         NA         NA

# find that the new algorithm does about 1000 times better than
#   the very first algorithm

```

# Conclusion

- ▶ The central message I wanted to convey: **algorithm matters**
- ▶ It's easy to boast the computer's hardware if you're rich
  - ▶ but it's more impressive if you beat everyone else with speed using the same hardware
  - ▶ (and practically, the only way out of a tech interview when algorithms are involved)
- ▶ Of course, this lecture note is more advanced than the rest and definitely won't appear in the final