

Functions

Yufeng Huang

Associate Professor of Marketing, Simon Business School

August 15, 2022

Samsung's smartphone quantity choice

- ▶ Samsung sells a generic smartphone and needs to determine how many units to produce
 - ▶ for now, think of Samsung as the only player in this market
 - ▶ “Samsung vs HTC” gives a richer case where HTC is a competing player
- ▶ Production cost is \$100 per unit
- ▶ Market determines price, with

$$P(q) = 500 - 0.6 \cdot q$$

- ▶ Question: how many units does Samsung produce maximize its profits?

Samsung's smartphone quantity choice

- ▶ Could hand-calculate this problem for a few number of quantities:
 - ▶ $q = 100, 200, 300$ and 400
 - ▶ for the above quantities, market prices are \$440, \$380, \$320, and \$260
 - ▶ price minus costs are \$340, \$280, \$220, and \$160
 - ▶ profits are \$34,000, \$56,000, \$66,000, and \$64,000
- ▶ \Rightarrow should produce 300 units

In code

```
# marginal costs
mc_S <- 100

# the four quantities
pos_qty <- c(100, 200, 300, 400)

# initialize the profit vector
profit_S <- rep(NA, length(pos_qty))

# loop to calculate the profits
for (q in 1:length(pos_qty)) {
  q_S <- pos_qty[q]
  price <- 500 - 0.6*q_S
  profit_S[q] <- q_S * (price - mc_S)
}

# examine results
data.frame(qty = pos_qty, profit = profit_S)

##   qty profit
## 1 100  34000
## 2 200  56000
## 3 300  66000
## 4 400  64000
```

What if market demand fluctuates?

- ▶ Market demand fluctuates: demand intercept is 200 in year 1, 300 in year 2, 400 in year 3

```
# Year 1:
profit_S_y1 <- rep(NA, length(pos_qty))

# loop to calculate the profits
for (q in 1:length(pos_qty)) {
  q_S <- pos_qty[q]
  price <- 200 - 0.6*q_S # note the intercept is 200
  profit_S_y1[q] <- q_S * (price - mc_S)
}

# examine results
data.frame(qty = pos_qty, profit = profit_S_y1)

##   qty profit
## 1 100  4000
## 2 200 -4000
## 3 300 -24000
## 4 400 -56000
```

```

# Year 3:
profit_S_y3 <- rep(NA, length(pos_qty))

# loop to calculate the profits
# you might see that we don't really need a loop to calculate profits,
# but we'll later extend this point to a separate point
for (q in 1:length(pos_qty)) {
  q_S <- pos_qty[q]
  price <- 400 - 0.6*q_S # note the intercept is 200
  profit_S_y3[q] <- q_S * (price - mc_S)
}

# examine results
data.frame(qty = pos_qty, profit = profit_S_y3)

##   qty profit
## 1 100  24000
## 2 200  36000
## 3 300  36000
## 4 400  24000

```

Alternative: write a function

```
# define a profit function
samsung_profit_vec <- function(intercept) {
  profit_S <- rep(NA, length(pos_qty))

  # loop to calculate the profits
  for (q in 1:length(pos_qty)) {
    q_S <- pos_qty[q]
    price <- intercept - 0.6*q_S    # note the intercept is an argument now
    profit_S[q] <- q_S * (price - mc_S)
  }

  # return
  return(profit_S)
}

# year 1
samsung_profit_vec(intercept = 200)    # can omit "intercept = "

## [1] 4000 -4000 -24000 -56000

# year 2
samsung_profit_vec(intercept = 300)

## [1] 14000 16000 6000 -16000

# year 3
samsung_profit_vec(intercept = 400)

## [1] 24000 36000 36000 24000
```

Alternative con'd: write a more flexible function if needed...


```
# define a profit function (more flexible)
samsung_profit_vec_2 <- function(intercept, pos_qty) {
  profit_S <- rep(NA, length(pos_qty)) # note: pos_qty is now an argument!

  # loop to calculate the profits
  for (q in 1:length(pos_qty)) {
    q_S <- pos_qty[q]
    price <- intercept - 0.6*q_S
    profit_S[q] <- q_S * (price - mc_S)
  }

  # return
  return(profit_S)
}
```



```
# ... so that we can plot profit at a range of different quantity vectors
plot(100:500, samsung_profit_vec_2(intercept = 400, pos_qty = 100:500),
     type = 'l', xlab = "quantity", ylab = "profit")
```



So, why do we need to write functions?

- ▶ We need to repeatedly call a routine
 - ▶ the routine *itself* is systematic
 - ▶ but not systematic for:
 - ▶ when we need it
 - ▶ how we use it
 - ▶ and more importantly there are no existing functions that achieve what we want to do
- ▶ In the toy example
 - ▶ how to compute Samsung's profit is the systematic rule
 - ▶ where do we apply the rule is more flexible (or *should be* more flexible??)
 - ▶ e.g., we can later write a function about Samsung's profit under any quantity q_S , instead of the vector of Samsung's profit under a given set of quantities

Anatomy of a function

Anatomy of a function

- ▶ A function has a name assigned to it
- ▶ Takes zero, one or several inputs known as *arguments*
- ▶ The expressions forming the operations comprise the body of the function

```
# built-in examples  
getwd() # no argument  
mean(x, na.rm = T) # one or two argument  
plot(y ~ x, data, type = 'l', lty = 2) # many arguments
```

- ▶ Returns a *single object*

Example: square

Write a function that squares its argument:

```
square <- function(x) {  
  return(x^2)  
}
```

- ▶ the function name is “square”
- ▶ the function has one argument
 - ▶ note: `square(x)` and `square(y)` are **calls** of the **same function**
- ▶ the function body consists of one simple expression “`x^2`”
- ▶ it returns the value “`x^2`”

Example: square

```
# works just like built-in functions in R
square(10)

## [1] 100

# works on vector too
square(1:5)

## [1] 1 4 9 16 25

# question for you: why??

# returns error when we put characters on it
square("Hi I'm Yufeng")

## Error in x^2: non-numeric argument to binary operator

# question for you: why??
```

Defining a function¹

The following are equivalent:

```
# standard
square <- function(x) {
  return(x^2)
}

# can skip return()
square <- function(x) {
  x^2
}

# can write simple expression in one line
square <- function(x) x^2

# [EASILY MISREAD!] can even split into multiple lines
square <- function(x)
  x^2
```

¹The last approach is very hard to read and easy to make mistake...

Side: Function name

- ▶ Cannot start with number or contain illegal characters

```
# all 3 are illegal
_square <- function(x) x^2

2power <- function(x) x^2

power-two <- function(x) x^2
```

- ▶ Can be defined as an operator
 - ▶ not advised for beginners because 1) harder to read and 2) can override built-in operators

```
# define the reverse of power operator
`%%` <- function(x, y) yx
3%2

## [1] 8

# much better for readability is
reverse_power <- function(power, base) basepower
# why? 1) base and power are more concrete to understand
#      2) function name tells the reader what it does
```


Only allow one output (“multiple outputs” => group as a list)

```
# square of sum: a function with two arguments
square_of_sum <- function(x, y) {
  return((x + y)^2)
}
square_of_sum(1, 2)

## [1] 9

# can ONLY return a single object
square_of_sum <- function(x, y) {
  x_y <- x + y
  x_y_2 <- x_y^2
  return(list(sum = x_y, sum_square = x_y_2))
}
square_of_sum(1, 2)$sum # first element in the list

## [1] 3

square_of_sum(1, 2)$sum_square # second element

## [1] 9
```

Arguments in a function

Can admit one or multiple arguments

```
# function with one argument
square(2)

## [1] 4

# function with two arguments
reverse_power <- function(power, base) base^power
reverse_power(3, 2)

## [1] 8

# function with multiple arguments
c(2, 1, 5, 2, 3, 6)

## [1] 2 1 5 2 3 6
```

Argument matching: can change order or arguments if they are clearly specified

```
# But, I don't quite remember what is what in reverse_power()

# can reverse argument explicitly
reverse_power(base = 2, power = 3)

## [1] 8

# cannot do so if we do not explicitly state input
reverse_power(2, 3)

## [1] 9
```

Argument matching: useful with functions with many (optional) arguments

```
# aggregate example
res_1 <- aggregate(x = mpg ~ cyl, data = mtcars, FUN = mean)

# I don't have to remember the order of stuff...
res_2 <- aggregate(data = mtcars, FUN = mean, x = mpg ~ cyl)

# Can partially match argument names (partial matching)
res_3 <- aggregate(mtcars, FUN = mean, mpg ~ cyl)

# and gives the identical result
identical(res_1, res_2)

## [1] TRUE

identical(res_2, res_3)

## [1] TRUE
```

Can admit no argument

```
# function with no argument
two_squared <- function() {
  x <- 2
  return(x^2)
}

# function with strictly no argument
two_squared()

## [1] 4

# returns error when you give it an argument (because it allows no argument)
two_squared(2)

## Error in two_squared(2): unused argument (2)
```

Arguments *can* have default values²

```
# function with default argument
squared <- function(x = 10) {
  return(x^2)
}

squared(2)

## [1] 4

squared()      # now gives default value x = 10

## [1] 100
```

²And this is different from no argument

Why have default values?

Default value is also very useful for functions with many arguments³

```
# plot, setting marker type as "points",  
# color as "black", and marker size as 1  
plot(mtcars$mpg, mtcars$wt, type = "p", col = "black", cex = 1)  
  
# I could just leave everything to default  
plot(mtcars$mpg, mtcars$wt)
```

³Results are the same but R recorded them a bit differently based on how some options are set

Can specify `missing()` to allow for missing arguments with no default

```
# sum over 3 numbers, one with default and one without
sum3 <- function(x, y, z = 0) {
  if (missing(y)) {
    res <- x + z
  } else {
    res <- x + y + z
  }
  return(res)
}

sum3(1)
## [1] 1

sum3(1, 2)
## [1] 3

sum3(1, 2, 3)
## [1] 6
```

Default value can depend on other variables

```
# default value as a function
x_plus_y <- function(x, y = log(x)) {
  return(x + y)
}

x_plus_y(2, 3)
## [1] 5


x_plus_y(2)      # y is missing! now take default y = log(x)
## [1] 2.693147

# WAIT: how did we know y = log(2)?
```

Can use dots (...) to pass more arguments to inside of a function

```
# Let's say we want to plot lines
lineplot <- function(x, ...) {
  plot(x, type = 'l', ...)
}

par(mfrow = c(1, 2))
# we can just plot a line
lineplot(1:10)
# or add other options
lineplot(1:10, xlab = "some index", ylab = "some values", main = "some title", col = "blue", lty = 2)
```



```
# but changing type will run into argument conflicts
lineplot(1:10, type = "p")
```

```
## Error in plot.default(x, type = "l", ...): formal argument "type" matched by multiple actual arguments
```

Built-in example

paste (base)

R Documentation

Concatenate Strings

Description

Concatenate vectors after converting to character.

Usage

```
paste (... , sep = " ", collapse = NULL)
paste0(... , collapse = NULL)
```

Arguments

... one or more R objects, to be converted to character vectors.
sep a character string to separate the terms. Not [NA_character_](#).
collapse an optional character string to separate the results. Not [NA_character_](#).

```
# paste()
paste("We", "are", "awesome")    # default for argument sep is " "

paste("We", "are", "awesome", "R", "Programmers")    # varies the number of arguments, still works

paste("We", "are", sep = " ", "awesome")    # argument matching; but this is confusing...

## [1] "We are awesome"
## [1] "We are awesome R Programmers"
## [1] "We are awesome"
```

Side: Operators as functions

Operators as functions

```
# `` is the notation to refer to an operator
`+`(2, 3)      # equivalent to 2 + 3

## [1] 5

`+`(2, `*(2, 3))

## [1] 8

2 + (2 * 3)    # equivalent

## [1] 8

# define an array
my.vec <- c("one", "two", "three", "four")      # 2 by 2 by 2 cube
`[`(my.vec, 2)

## [1] "two"

my.vec[2]      # equivalent

## [1] "two"
```

A more advanced example: define your own operators⁴


```
# Courtesy of Hadley Wickham, let's do a prank
#  runif() generates random number between 0 and 1
#  evenly distributed so 15% of runif(length(x)) will fall below 0.15

# now let's define a function `(` -- as the parenthesis operator
`(` <- function(x) {
  ifelse(
    runif(length(x)) < 0.15, # condition
    x + 2, # if TRUE
    x      # if FALSE
  )
}
```

⁴Strictly speaking, `runif()` funds uniform random numbers

A more advanced example: define your own operators

```
# now, plotting (1:100) will not work because `(` has effect on 1:100  
plot((1:100)) # note that plot(1:100) does not change
```




```
# remove this to be safe  
rm(`(`)
```


What we know and what we don't know

- ▶ What we know
 - ▶ `fun_name <- function(x) { body }`
 - ▶ argument and argument matching
- ▶ What we don't know
 - ▶ lazy evaluation
 - ▶ environment and scoping
- ▶ And importantly: more applications

R is lazy



Lazy evaluation 1: only needs an argument when actually evaluating it⁵

```
# Consider a function with arguments a and b
g <- function(a, b) {
  return(a * a * a)
}

# call g(2) will be equivalent to calling g(2, 1)
g(2, 1) # standard call, but b is useless

## [1] 8

g(2)    # ignores b

## [1] 8
```

⁵Lazy Evaluation IS UNIQUE in R!!

Lazy evaluation 2: runs into error when – not before – calling arguments that do not exist

```
# Consider a function with arguments a and b
g <- function(a, b) {
  print(a * a * a)
  print(b)
}

# now g(2) will be different from g(2, 1)
g(2, 1) # standard call, now b is useful

## [1] 8
## [1] 1

g(2)    # still evaluates print(a^3)

## [1] 8

## Error in print(b): argument "b" is missing, with no default
```

Lazy evaluation 3: default arguments evaluated only when called upon

```
# Consider a function with arguments a and b
g <- function(a, b = 3*a) {
  print(a * a * a)
  a <- a^3
  print(b)
}

# now g(2, 1) takes standard values given by arguments
g(2, 1)

## [1] 8
## [1] 1

# but g(2) uses b that comes from the UPDATED a
g(2)      # a updated to 2^3, then default value of b is 3*8

## [1] 8
## [1] 24
```

Environments and variable scoping

Motivating example: what is x?

```
# What is x after the following?

rm('x') # remove any x in the memory

## Warning in rm("x"):  object 'x' not found

f <- function(x) {
  x <- x^2
  return(x)
  x^2
}
f(2)

## [1] 4

# NOW WHAT IS x?
```

```
x

## Error in eval(expr, envir, enclos):  object 'x' not found

# x does not exist
# WHY?
```

Variables in local vs global environments

```
# How does this structure work?
y <- 10
f <- function(x) {
  return(x + y)
}
f(0)

## [1] 10

# How does this different structure work?
y <- 10
g <- function(x) {
  y <- 5
  return(x + y)
}
g(0)

## [1] 5

# After running these lines, what's y?
y

## [1] 10
```



```
# Can print the y variables in given environment to see what happens...
```

```
# How does this structure work?
```

```
y <- 10
f <- function(x) {
  print(paste("y is", y)) # local y
  return(x + y)
}
f(0)
```

```
## [1] "y is 10"
```

```
## [1] 10
```

```
# How does this different structure work?
```

```
y <- 10
g <- function(x) {
  y <- 5
  print(paste("y is", y)) # local y
  return(x + y)
}
g(0)
```

```
## [1] "y is 5"
```

```
## [1] 5
```

```
# After running these lines, what's y?
```

```
print(paste("y is", y)) # global y
```

```
## [1] "y is 10"
```

What happened?

- ▶ In the first example, can't find local `y` so search the global environment
- ▶ In the second example, `y` is defined in local environment so use local `y`
- ▶ While running `g()`, global `y` is not changed, so `y = 5` is only for inside `g()`
- ▶ Therefore, after running `g()`, `y` in the global environment remains unchanged

Scope of a variable

```
# The scope of a variable is the range of places where you can find it
# R will first try to find the variable in the current environment
y <- 10
paste("y in global is", y)      # global

## [1] "y in global is 10"

f <- function() {
  y <- 5
  print(paste("y in f() is", y))
}
f()      # local to function f

## [1] "y in f() is 5"
```

Scope of a variable

```
# The scope of a variable is the range of places where you can find it
# R will first try to find the variable in the current environment

# If there is no such variable y, R will go up and find in parent environment
f <- function() {
  g <- function() {
    print(paste("y in g() is", y))
    # no y in g, search in environment of f
  }
  print(paste("y in f() is", y)) # no y; search global environment
  y <- 5
  return(g())
}
f()      # first the global y and then the local-to-f() y

## [1] "y in f() is 10"
## [1] "y in g() is 5"
```

Your turn: scope of a variable

```
# clear all previous memory
rm(list = ls())

# g() finds y in the environment of f()
f <- function(x) {
  y <- 10
  g <- function(x) {
    x + y
  }
  return(g(x))
}
f(5)

## [1] 15

# how about this?
f <- function(x) {
  y <- 10
  return(g(x))
}
g <- function(x) {
  return(x + y)
}
# f(5) is ?
```

What happened?

- ▶ In the first example, `g()` is defined in the environment of `f()`
 - ▶ `f()` passes `x` to `g()`
 - ▶ `g()` needs `y`, so searches the parent environment (i.e., that of `f()`) to find `y`
 - ▶ `y` is defined in the environment of `f()`
- ▶ In the second example, `f()` is defined in the global environment
 - ▶ `f()` calls `g()` inside, can't find `g()` inside its own environment
 - ▶ finds `g()` in the parent environment, so global
 - ▶ `g()` is defined in global so calls `y` in global
 - ▶ can't find it, returns error

A more advanced example

```
# clear all previous memory
rm(list = ls())

# we first assign a function of which the OUTPUT IS A FUNCTION
j <- function(x) {
  y <- 2
  return(function() c(x, y))      # returns this function
}
k <- j(1)      # so k is a function

# side: this is different from
# j <- function(x) {
#   y <- 2
#   return(c(x, y))
# }

# note that k maintains the environment within j
#   so even call k in global
k()

## [1] 1 2

#   it gives the y value in j()
```

Examples / exercises

Example 1: check positive value

```
# Let's write a function that checks positive value
check_positive <- function(x) {
  if (x > 0) print(TRUE) else print(FALSE)
}

# instead of printing some value, more often we "return" the value
check_positive <- function(x) {
  return(x > 0)
}
```

Example 1: check positive value

```
# simple check_positive
check_positive <- function(x) {
  return(x > 0)
}
```

```
# Your turn: output?
```

```
check_positive(0)
```

```
check_positive(NA)
```

```
check_positive(TRUE)
```

```
check_positive("positive")
```

```
check_positive(-2:2)
```

Example 2: root of quadratic equations⁶

- ▶ Suppose I want to evaluate the root x of quadratic equations with the general form

$$a \cdot x^2 + b \cdot x + c = 0$$

where a , b and c are parameters

- ▶ Mathematically there are two solutions of x , jointly expressed as

$$x = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$$

⁶Which we talked about in Week 1

Example 2: root of quadratic equations

```
quadratic_root <- function(a, b, c) {  
  sqr_root <- sqrt(b^2 - 4*a*c)  
  x1 <- (-b + sqr_root) / (2*a)  
  x2 <- (-b - sqr_root) / (2*a)  
  return(list(root1 = x1, root2 = x2))  
}
```

```
# root of  $-x^2 + 2x + 1 = 0$   
x <- quadratic_root(-1, 2, 1)  
x$root1  
  
## [1] -0.4142136  
  
x$root2  
  
## [1] 2.414214
```

```
# root of  $x^2 + x + 1 = 0$ 
x <- quadratic_root(1, 1, 1)

## Warning in sqrt(b^2 - 4 * a * c): NaNs produced

# warning and no result in real number (returns NaN)
x$root1

## [1] NaN
```

```

# instead, wrap around a warning message and return complex number
quadratic_root <- function(a, b, c) {
  if (b^2 - 4*a*c >= 0) {
    sqr_root <- sqrt(b^2 - 4*a*c)
  } else {
    warning("equation has no real root, complex root returned")
    sqr_root <- sqrt(- b^2 + 4*a*c)*1i
  }
  x1 <- (-b + sqr_root) / (2*a)
  x2 <- (-b - sqr_root) / (2*a)
  return(list(x1, x2))
}
quadratic_root(1, 1, 1)

## Warning in quadratic_root(1, 1, 1): equation has no real root,
## complex root returned

## [[1]]
## [1] -0.5+0.8660254i
##
## [[2]]
## [1] -0.5-0.8660254i

```

The warning() and stop() functions produce warning and error messages

```
gen_error <- function(x) {  
  stop("hey this is an error code and function stops here")  
  return(x^2)  
}  
gen_error(2)  
  
## Error in gen_error(2): hey this is an error code and function stops  
here  
  
gen_warning <- function(x) {  
  warning("hey this is a warning message and the code goes on")  
  return(x^2)  
}  
gen_warning(2) # does not stop the code!  
  
## Warning in gen_warning(2): hey this is a warning message and the  
code goes on  
  
## [1] 4
```

Example 3 and optimization

Example 3: profit maximization

- ▶ Recall: Samsung wants to maximize profit, which depends on its quantity, demand level (how good the market is), and marginal costs

$$\Pi(q) = q \times (P(q) - c)$$

- ▶ Inverse demand is

$$P(q) = \text{demand} - 0.6 \times q$$

- ▶ We want to know how to set prices that maximizes Π_t given the demand intercept (“demand”) and cost

Let's write the profit function

```
# demand function
# note: careful to include variables as arguments (prevent scoping)
market_price <- function(quantity, intercept) {
  price <- intercept - 0.6 * quantity
  return(price)
}



# profit function
profit <- function(quantity, intercept, cost) {
  price <- market_price(quantity, intercept)
  profit <- quantity*(price - cost)
  return(profit)
}
```

In other words, `profit()` is a function that nests the function `market_price()`

Why nesting functions?

```
# why nest function? Mostly because
#   1) I need market_price() somewhere else or
#   2) I need to call market_price() multiple times (e.g. this plot)
pos_qty <- 100:500          # possible quantities on the x-axis
par(mfrow = c(1, 2))       # allocate left/right graph



# demand under different market conditions
for (incpt in c(250, 400)) {
  plot(pos_qty, market_price(pos_qty, incpt),
       type = 'l', ylim = c(0, 350), ylab = "market price")
}
```



Now let's optimize Samsung's quantity across different demand conditions

##	year	demand_intercept	cost
## 1	1	200	93.357
## 2	2	300	87.213
## 3	3	200	86.581
## 4	4	400	93.752
## 5	5	500	82.383
## 6	6	400	85.532
## 7	7	300	71.438
## 8	8	200	90.128
## 9	9	300	105.996
## 10	10	400	81.100

Demand and costs in 10 years




How would we maximize profit? Brute-force approach

- ▶ A brute-force approach is to solve for all profits under all possible quantities
- ▶ Question: what's the problem with this approach?

```
# brute-force approach: take all possible quantities for each year
pos_qty <- 50:600

# example: year 1
profit_vec_1 <- profit(pos_qty, df$demand_intercept[1], df$cost[1])

# plot
plot(pos_qty, profit_vec_1,
      xlab = "quantity", ylab = "profit in Year 1", type = 'l')
```



```
# which quantity solves the highest profit?
pos_qty[which.max(profit_vec_1)]

## [1] 89
```


How would we maximize profit if this is Excel?

Solver Parameters

Set Objective:

To: ☒ Max ☐ Min ☐ Value Of:

By Changing Variable Cells:

Subject to the Constraints:

Add

Change

Delete

Reset All

Load/Save

☒ Make Unconstrained Variables Non-Negative

Select a Solving Method:

Solving Method

Select the GRG Nonlinear engine for Solver Problems that are smooth nonlinear. Select the LP Simplex engine for linear Solver Problems, and select the Evolutionary engine for Solver problems that are non-smooth.

Help Solve Close

A better approach: optim()

```
# optim() is like an Excel solver, but in R (and much better than Excel)
optimization_result <- optim(
  par = 100,          # initial trial value
  fn = function(qty) {
    - profit(qty, df$demand_intercept[1], df$cost[1])
  }                  # function with ONE ARGUMENT to MINIMIZE
)

# returns a list of results
optimization_result[1:2]

# first is quantity, second is -profit

## $par
## [1] 88.86719
##
## $value
## [1] -4738.68
```

Repeat this for many years

```
# for loop to wrap this optimization routine
df$optimal_quantity <- rep(NA, 10)
for (i in 1:nrow(df)) {
  df$optimal_quantity[i] <- optim(par = 100, function(qty) {
    - profit(qty, df$demand_intercept[i], df$cost[i])
  }) # the function will look for i outside of its body
}
df$market_price <- market_price(df$optimal_quantity, df$demand_intercept)

head(round(df, 3), n = 10)
```

##	year	demand_intercept	cost	optimal_quantity	market_price
## 1	1	200	93.357	88.867	146.680
## 2	2	300	87.213	177.305	193.617
## 3	3	200	86.581	94.512	143.293
## 4	4	400	93.752	255.215	246.871
## 5	5	500	82.383	348.008	291.195
## 6	6	400	85.532	262.070	242.758
## 7	7	300	71.438	190.312	185.812
## 8	8	200	90.128	91.562	145.062
## 9	9	300	105.996	161.660	203.004
## 10	10	400	81.100	265.742	240.555

Summary

- ▶ **Assigning** a function
 - ▶ `function_name <- function(argument1, argument2) {
 body_of_the_function }`
- ▶ **Calling** a function
 - ▶ `function_name(argument1, argument2)`
- ▶ Properties of the function (important things about arguments)
 - ▶ “...” passes arguments
 - ▶ argument matching
 - ▶ lazy evaluation
 - ▶ variable scoping
- ▶ `optim()`