

Recency, Frequency and Monetary Value

Yufeng Huang

Associate Professor of Marketing, Simon Business School

August 18, 2022

RFM targeting

- ▶ You already know the meaning of segmentation from your econ/marketing class
- ▶ RFM is a classical way to segment the customers and target the valuable sub-groups
 - ▶ internal logic: human behavior is persistent and therefore past choices of a consumer can predict future choices
 - ▶ not the ideal targeting approach but this is something easy to compute
- ▶ Of course, my real purpose is for you to practice for-loops

Prep-steps

```
# ===== CLEAR EVERYTHING =====
rm(list = ls())

# ===== READ TRIAL DATA =====
url <- 'https://dl.dropboxusercontent.com/s/xxfloksp0968mgu/CDNOW_sample.txt'
if (!file.exists('CDNOW_sample.txt')) download.file(url, 'CDNOW_sample.txt')
df.raw <- read.fwf('CDNOW_sample.txt', width = c(6, 5, 9, 3, 8), stringsAsFactors = F) # load data

# a) generate year and month (using date functions, but you can use all sorts of methods)
df.raw[[1]] <- NULL # drop old id
names(df.raw) <- c("id", "date", "qty", "expd")
df.raw$date <- as.Date(as.character(df.raw$date), "%Y%m%d")
df.raw$year <- as.numeric(format(df.raw$date, '%Y'))
df.raw$month <- as.numeric(format(df.raw$date, '%m'))

# b) aggregate into monthly data with number of trips and total expenditure
df.raw$trips <- 1
df.month <- aggregate(cbind(qty, expd, trips) ~ id + year + month, data = df.raw, FUN = sum)
# note 1: cbind() ~ ... will aggregate using the same function, but on multiple variables

# c) generate a table of year-months, merge, replace no trip to zero.
id.year.month <- expand.grid(list(id = 1:max(df.raw$id),
                                year = unique(df.raw$year), month = unique(df.raw$month)))
id.year.month$t <- (id.year.month$year - 1997)*12 + id.year.month$month
df <- merge(id.year.month, df.month, by = c("id", "year", "month"), all = TRUE)
df <- df[df$year == 1997 | df$month <= 6, ] # subset, get rid of 199807-199812
df <- df[order(df$id, df$year, df$month), ] # sort
df[is.na(df$qty), c("qty", "expd", "trips")] <- 0 # replace no-trip months to zero
```

Recency (1)

```
# Method 1
# let's keep track of time
t0 <- proc.time()      # log current system time

# initialize recency
df$recency <- NA

# double for-loop, on individual and then year-month
# note: indentation for readability
for (i in 1:max(df$id)) {
  df$recency[df$id == i & df$t == 1] <- NA
  for (m in 2:18) {
    if (df$qty[df$id == i & df$t == m - 1] > 0) {
      df$recency[df$id == i & df$t == m] <- 1
    } else {
      df$recency[df$id == i & df$t == m] <- df$recency[df$id == i & df$t == m - 1] + 1
    }
  }
}

t1 <- proc.time()

# confirm that it works
head(df, 4)

##   id year month t qty  expd trips recency
## 1  1 1997     1 1   4 59.06      2      NA
## 2  1 1997     2 2   0  0.00      0       1
## 3  1 1997     3 3   0  0.00      0       2
## 4  1 1997     4 4   0  0.00      0       3

# elapsed time
t1 - t0

##   user  system elapsed
## 4.86   2.00   6.86
```

Thoughts?

- ▶ Some of you did year and then month as two separate loop (so in total 3 layers of loops)
 - ▶ at first glance straightforward, because the variables you have are “year” and “month” separate
 - ▶ however, does not add value to the *structure*, because the logic of the algorithm should be
 - ▶ “continue to the next month, for the same individual”, rather than
 - ▶ “continue to the next month, for the same individual in the given year”
 - ▶ therefore, it is much easier to construct an additional variable, here \$t
- ▶ Still, we have two layers of loops, and they take quite some time
 - ▶ any ways to improve the code?

Thoughts?

- ▶ Note that the only reason we need individual i is that we try to treat the first period separately
 - ▶ if it is the first month for a given i , we reset value for recency to NA
 - ▶ for second month onwards, we let recency continue from its last value
- ▶ Could write the code into a single-layer loop
 - ▶ now, iterator is “an observation” rather than “an individual”
 - ▶ at each step, check whether we are now looking at a new individual or not
 - ▶ if yes, reset recency to NA
 - ▶ if no, let recency continue from its last value

```

# Method 2
t0 <- proc.time()

# sort df (should be sorted, just to make sure)
df <- df[order(df$id, df$t), ]

# initialize recency
df$recency <- NA

# single for-loop, on observation number (note that we start from 2)
for (obs in 2:nrow(df)) {
  if (df$id[obs] != df$id[obs - 1]) { # check whether id switched
    df$recency[obs] <- NA             # if id switched, assign NA to current obs of data
  } else {
    if (df$qty[obs - 1] > 0) {
      df$recency[obs] <- 1           # if bought last month, reset recency to 1
    } else {
      df$recency[obs] <- df$recency[obs - 1] + 1      # if not, advance recency by 1
    }
  }
}

t1 <- proc.time()

# works exactly the same way
head(df, 4)

##   id year month t qty  expd trips recency
## 1  1 1997     1 1   4 59.06      2      NA
## 2  1 1997     2 2   0  0.00      0       1
## 3  1 1997     3 3   0  0.00      0       2
## 4  1 1997     4 4   0  0.00      0       3

# 15-20 times faster (!!!) because we only have one-layer of loop
t1 - t0

##   user  system elapsed
## 0.27   0.09   0.36

```

Frequency (1)

```
# One might think about writing a loop for frequency
#   but note that frequency is just the last quarter's total number of trips
#   now we realize that we do not need loops for this

# construct quarter
df$quarter <- (df$t - 1) %/% 3 + 1 # 6 quarters in total

# aggregate to get frequency
frequency <- aggregate(trips ~ id + quarter, df, sum) # number of trips in a quarter
frequency$quarter <- frequency$quarter + 1 # so it later merges with next quarter in df
colnames(frequency)[3] <- "frequency"

# merge back to data
df <- merge(df, frequency, by = c("id", "quarter"), all.x = TRUE)

# observe
head(df, 4)
```

##	id	quarter	year	month	t	qty	expd	trips	recency	frequency
## 1	1	1	1997	1	1	4	59.06	2	NA	NA
## 2	1	1	1997	2	2	0	0.00	0	1	NA
## 3	1	1	1997	3	3	0	0.00	0	2	NA
## 4	1	2	1997	4	4	0	0.00	0	3	2

Frequency (2)

```
# We can of course get it from a loop

# initialize frequency
df$frequency <- NA

# for-loop to get frequency
for (i in 1:max(df$id)) {
  for (q in 2:6) {
    df$frequency[df$id == i & df$quarter == q] <-
      sum(df$trips[df$id == i & df$quarter == q - 1])
  }
}

# observe
head(df, 4)
```

##	id	quarter	year	month	t	qty	expd	trips	recency	frequency
## 1	1	1	1997	1	1	4	59.06	2	NA	NA
## 2	1	1	1997	2	2	0	0.00	0	1	NA
## 3	1	1	1997	3	3	0	0.00	0	2	NA
## 4	1	2	1997	4	4	0	0.00	0	3	2

Monetary value (1)

```
# Idea: average non-zero monthly expenditure in the past
# so we need: 1) average of nonzero months
#           2) do so by different sets of time periods (first to previous month for each id)

# Method 1
t0 <- proc.time()
df$monvalue <- NA
df <- df[order(df$id, df$t), ]

# directly loop over rows
for (obs in 1:nrow(df)) {
  if (df$t[obs] == 1) { # if this is initial observation
    expd.temp <- numeric(0) # initialize an empty vector if no history yet
  } else {
    expd.temp <- c(expd.temp, df$expd[obs - 1]) # expand vector of expenditure
    expd.temp[expd.temp == 0] <- NA # replace zeros into NAs
    df$monvalue[obs] <- mean(expd.temp, na.rm = T) # take mean
  }
}

t1 <- proc.time()

# check result
head(df, 4)

##   id quarter year month t qty  expd trips recency frequency monvalue
## 1  1         1 1997     1 1    4 59.06      2     NA         NA      NA
## 2  1         1 1997     2 2    0  0.00      0      1         NA  59.06
## 3  1         1 1997     3 3    0  0.00      0      2         NA  59.06
## 4  1         2 1997     4 4    0  0.00      0      3          2  59.06

t1 - t0

##   user  system elapsed
##   0.31    0.06    0.37
```

Monetary value (2)

```
# Idea: we can just keep track of scalar values
#   realizing that average monthly expenditure is total expd / number of months

# Variant 1: avoid vector expansion
t0 <- proc.time()
df$monvalue <- NA
df <- df[order(df$id, df$t), ]

# this time, keep sum and total number as two scalar variables
for (obs in 1:nrow(df)) {
  if (df$t[obs] == 1) {
    sum.expd <- 0                                # initialize their value at the start of each individual
    total.months <- 0
  } else {
    if (df$trips[obs - 1] > 0) {                  # if there are some trips
      sum.expd <- sum.expd + df$expd[obs - 1]
      total.months <- total.months + 1
    }
    df$monvalue[obs] <- sum.expd / total.months
  }
}

t1 <- proc.time()

# takes roughly 10-20% less time because we don't have vector expansion
t1 - t0

##      user  system elapsed
##    0.27    0.00    0.26
```

Monetary value (3)

```
# Idea: can use an index to keep track of the initial obs for each individual

# Variant 2: avoid defining a local vector
t0 <- proc.time()
df$monvalue <- NA
df <- df[order(df$id, df$t), ]

# turn zero expenditure into NA
df$expd[df$expd == 0] <- NA

# this time, subset "on the fly" and make use of na.rm = T
for (obs in 1:nrow(df)) {
  if (df$t[obs] == 1) {
    init <- obs      # initial observation for the given individual
  } else {
    df$monvalue[obs] <- mean(df$expd[init:(obs - 1)], na.rm = T)  # subset, but do not copy
  }
}

t1 <- proc.time()

df$expd[is.na(df$expd)] <- 0    # return the zeros

# no time saving, presumably because mean() sums up a range of values
# and there're wasted calculations across time periods
t1 - t0

##      user  system elapsed
##    0.34    0.03    0.38
```

Thoughts?

- ▶ Looping over observation index rather than looping over individual-time (worse is individual-year-month in 3 layers) again saves a lot of time
- ▶ In addition, in the baseline method here, some time is wasted when we dynamically expand the vector `expd.temp`
 - ▶ does not help when we subset and take averages because of wasteful calculations in the `mean(vector)`
 - ▶ efficiency can be improved by observing that $\text{mean} = \text{sum} / \text{nr. months}$, and keeping track of scalar sums and number is more efficient

Finally, RFM index

```
# compute RFM index
df$index <- -0.1*df$recency + 3.5*df$frequency + 0.2*df$monvalue

# split into quantiles
# note: cut() itself will generate a factor, in this case happened to be nicely ordered
# note: perfectly fine to use a for-loop for this, back to this in the apply lecture notes
df$quantiles <- as.numeric(cut(df$index, quantile(df$index, seq(0, 1, 0.1), na.rm = T)))
# show parts of data
df[1:18, c(1, 5, 7, 8, 9, 10, 11, 12, 13)]
```


##	id	t	expd	trips	recency	frequency	monvalue	index	quantiles
## 1	1	1	59.06	2	NA	NA	NA	NA	NA
## 2	1	2	0.00	0	1	NA	59.06	NA	NA
## 3	1	3	0.00	0	2	NA	59.06	NA	NA
## 4	1	4	0.00	0	3	2	59.06	18.512	9
## 5	1	5	0.00	0	4	2	59.06	18.412	9
## 6	1	6	0.00	0	5	2	59.06	18.312	9
## 7	1	7	0.00	0	6	0	59.06	11.212	8
## 9	1	8	14.96	1	7	0	59.06	11.112	8
## 8	1	9	0.00	0	1	0	37.01	7.302	6
## 10	1	10	0.00	0	2	1	37.01	10.702	8
## 11	1	11	0.00	0	3	1	37.01	10.602	8
## 12	1	12	26.48	1	4	1	37.01	10.502	8
## 14	1	13	0.00	0	1	1	33.50	10.100	8
## 15	1	14	0.00	0	2	1	33.50	10.000	8
## 13	1	15	0.00	0	3	1	33.50	9.900	8
## 17	1	16	0.00	0	4	0	33.50	6.300	6
## 16	1	17	0.00	0	5	0	33.50	6.200	6
## 18	1	18	0.00	0	6	0	33.50	6.100	5

Alternatives to the cut() function?

```
# can write a loop over the 10 possible groups
qtile_cutoffs <- quantile(df$index, seq(0.1, 1, 0.1), na.rm = T)

df$quantiles <- 1
for (i in 2:10) {
  df$quantiles <- ifelse(df$index > qtile_cutoffs[i-1] &
                        df$index <= qtile_cutoffs[i], # test: if between two neighboring cutoffs
                        i, # yes: take group i
                        df$quantiles) # no: unchanged
}
```

```
# per-month expenditure
expd.by.group <- aggregate(expd ~ quantiles, df, mean)
# bar plot
barplot(expd.by.group[, 2], names.arg = 1:10,
        ylab = "average monthly expenditure", xlab = "deciles in the RFM index")
```




Is RFM useful at targeting consumers?

- ▶ Recall that we only use historical data to generate RFM
 - ▶ that means, no guarantee that high-RFM consumers will buy today
 - ▶ but we see that they do spend dramatically more
- ▶ Therefore, yes, RFM is useful at targeting consumers
 - ▶ why? Consumer behavior is persistent because different people have different preference/habit/purchase patterns
 - ▶ and we should target from top to bottom deciles depending on the marketing cost

Anyone tried to plot per-trip expenditure?

```
# per-TRIP expenditure (for existing trips)
trip.expd.by.group <- aggregate(expd/trips ~ quantiles, df, mean)
# bar plot
barplot(trip.expd.by.group[, 2], names.arg = 1:10,
        ylab = "average per-trip expenditure", xlab = "deciles in the RFM index")
```



Why are there such differences?

- ▶ Conditional on coming to the store, expenditure is much higher
 - ▶ (mechanically) because many individual-months have zero expenditure
- ▶ Do RFM predict expenditure amount, given that they are positive?
 - ▶ well, yes, but to a lesser extent compared to predicting *whether* expenditure is non-zero
- ▶ I.e., the index we constructed is good at predicting who comes when, but not how much they spend given that they come
 - ▶ these are both good marketing insights and should be examined separately